

# MEMORY-AUGMENTED RESERVOIR COMPUTING



## Abstract

This seminar explores implementing **external working memories** to reservoir models, with a focus on theoretical aspects of computational expressiveness.

The landscape of recent memory-augmented reservoir computing is examined, showing the Reservoir memory machine (**RMM**) as well as the Reservoir stack machine (**RSM**).

State of the art **end-to-end models** are then briefly investigated, along with some interesting future research directions.

# MOTIVATION



## Neural networks with external memories

Neural Turing Machines<sup>1</sup> and Differentiable Neural Computers<sup>2</sup> have been explored extensively within the ML landscape.

The main motivation behind such systems is to enable the network to model **long term dependencies**, for which standard RNNs tend to fail.

1. [Alex Graves, Greg Wayne, Ivo Danihelka \(2014\)](#)
2. [Alex Graves, et al. \(2016\)](#)

# MOTIVATION

Being notoriously hard to train in an **end-to-end** fashion (as in deep learning), the application of such networks is currently limited.

Here we will focus in particular on **memory-augmented Reservoir Computing** methods, which have been showing promising results lately.



## Advantages

Some advantages of these randomized recurrent architectures include much **lower training and inference times**, compared to their deep counterparts, along with some nice theoretical results.

# MOTIVATION

## Limitations

While proving useful in some scenarios, memory-augmented reservoir computing still hold onto some problems, like the need to produce **target memory control sequences** for the training algorithm. These aspects will be analyzed later on.

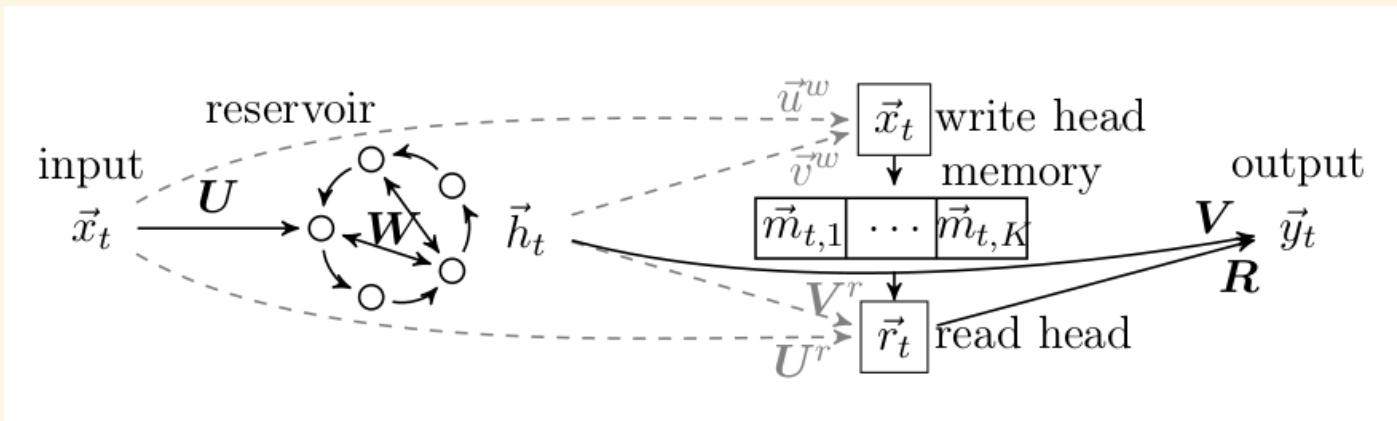
# IDEA



## Extending ESNs

The main idea behind memory-augmented reservoir computing, is that of extending Echo State Networks (ESN) with external memories.

By separating memory from computation we can thus design systems that are able to model arbitrarily long term dependencies in the input without interference<sup>1</sup>.



1. Benjamin Paassen, Alexander Schulz (2020)

# ECHO STATE NETWORKS

Echo State Networks are a kind of Randomized RNNs and are studied under the field of **Reservoir Computing**.

A standard echo state network follows the same update equations of a vanilla RNN:



## Definition (RNN equations)

$$\begin{aligned} h_t &= f(x_t, h_{t-1}) = \tanh(\mathbf{U} \cdot x_t + \mathbf{W} \cdot h_{t-1}) \\ y_t &= \mathbf{V} \cdot h_t \end{aligned}$$

With the following differences:

1.  **$\mathbf{U}$  and  $\mathbf{W}$  are fixed**
2. **The network must hold the Echo State Property.**

# ECHO STATE PROPERTY

Provides a theoretical that the state of the reservoir will depend on the input and not on the **initial conditions**, with past influences degrading over time. This is related to the concept of **Markovianity**<sup>1</sup>, which is typically achieved by designing the model to be a contractive map.



## Definition (*Contractive reservoir*)

Let  $\mathcal{R} = (\mathbf{U}, \mathbf{W}, b, \sigma, h_0)$  be a reservoir. We say  $\mathcal{R}$  is **contractive** with constant  $C \in (0, 1)$  if for any  $h, h' \in \mathbb{R}^m$  and any  $x \in \mathbb{R}^n$  it holds:

$$||f(x, h) - f(x, h')|| \leq C \cdot ||h - h'||$$

where  $f$  is defined as in the update equations of a vanilla RNN.

1. Claudio Gallicchio, Alessio Micheli (2011)

# LIMITATIONS OF ESN



## Memory capacity

Since echo state network hold the echo state property, the influence of past inputs is limited by the contractive properties of the network. As such, **long term dependency** tasks are out of reach of such models.

The memory capacity is limited by the number of neurons of the reservoir<sup>1</sup>

1. Igor Farkaš, Radomír Bosák, Peter Gergel' (2016)

# RESERVOIR MEMORY MACHINE



## Echo state property

The echo state property is necessary in order for the network not to discriminate on initial conditions, but it forcefully limits its capability to model long term dependencies in the input.

For this reason, **Reservoir memory machines** have been proposed.

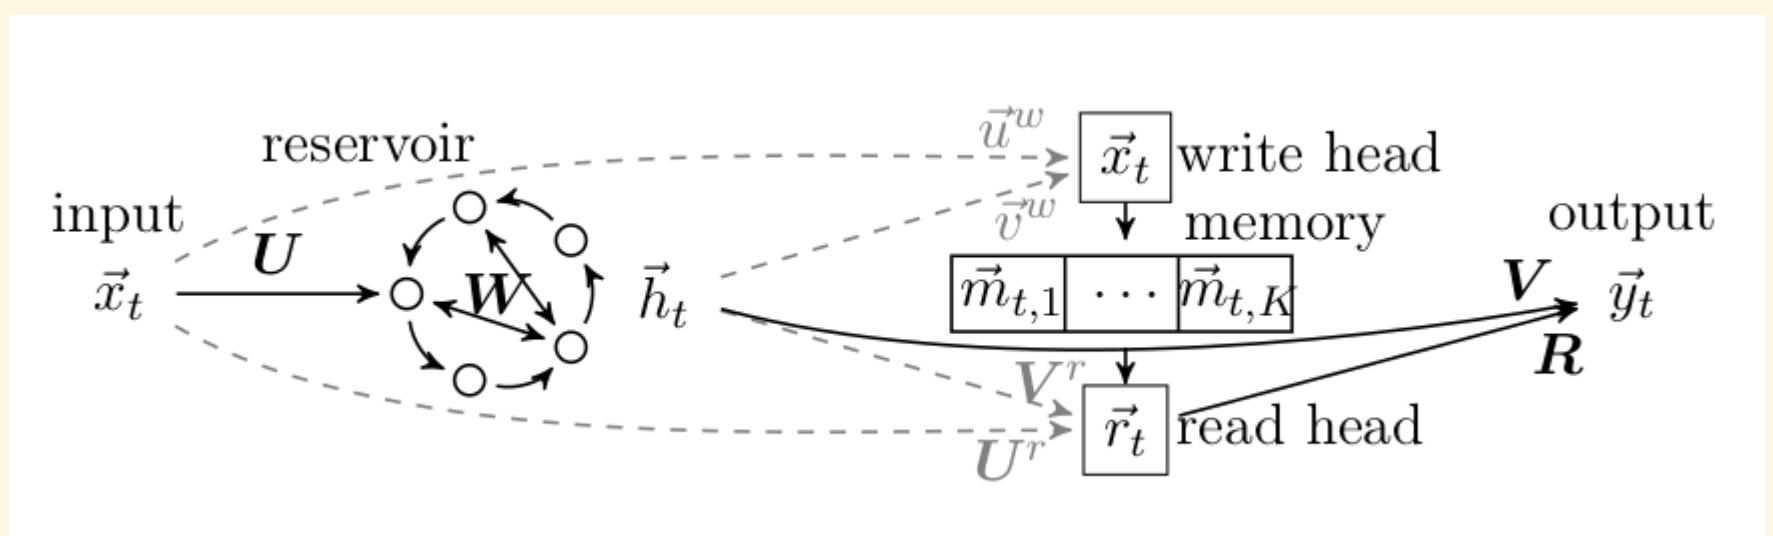
# RESERVOIR MEMORY MACHINE

The vanilla RMM<sup>1</sup> has been introduced as an extension of standard ESNs. In particular the network is equipped with an external memory that it can access using a **read** and a **write** head.

1. [Benjamin Paassen, Alexander Schulz \(2020\)](#)

# RESERVOIR MEMORY MACHINE

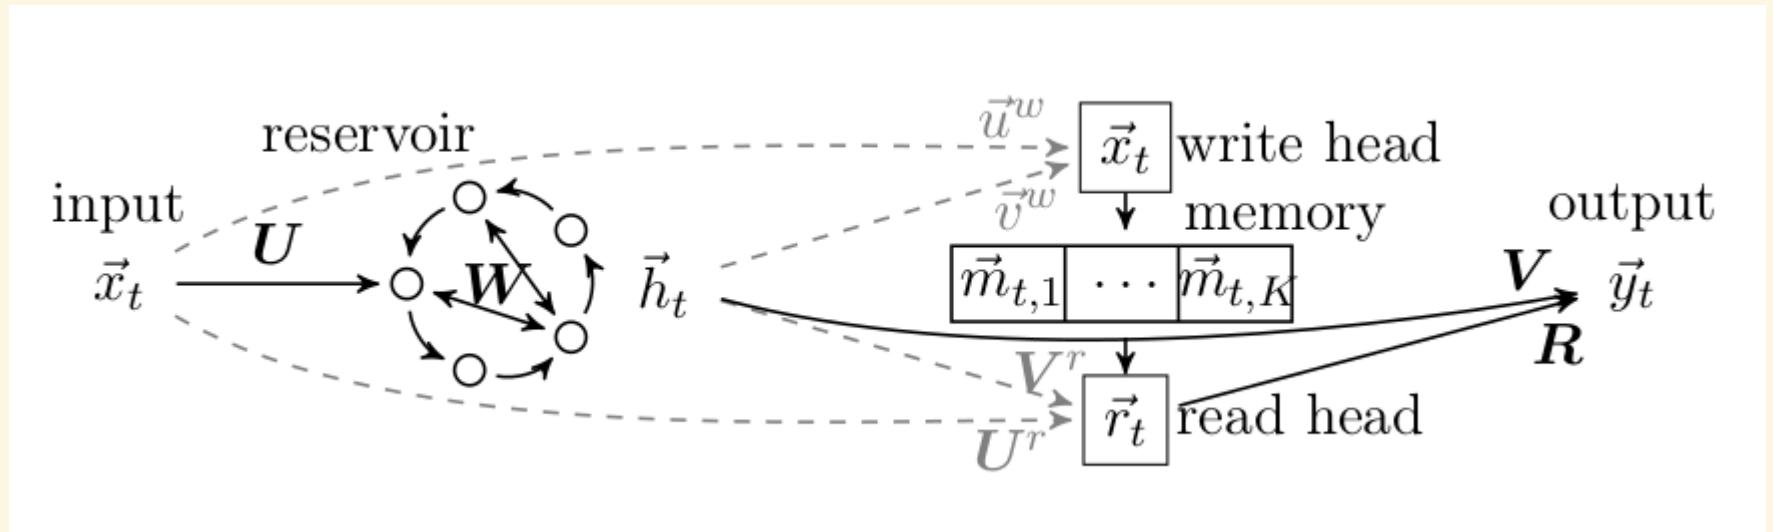
The vanilla RMM<sup>1</sup> has been introduced as an extension of standard ESNs. In particular the network is equipped with an external memory that it can access using a **read** and a **write** head.



1. Benjamin Paassen, Alexander Schulz (2020)

# RESERVOIR MEMORY MACHINE

## DYNAMICS



1. Copy previous memory state:  $\mathbf{M}_t \leftarrow \mathbf{M}_{t-1}$
2. Compute write head discriminator:  $c_t^w = u^w \cdot x^t + v^w \cdot h_t$ 
  1. if  $c_t^w > 0$  then  $\mathbf{M}_{t,k} \leftarrow x_t$  and  $k_t \leftarrow k_{t-1} + 1$
3. Compute read head discriminator:  $c_t^r = \mathbf{U}^r \cdot x_t + \mathbf{V}^r \cdot h_t$ 

$l_t$	if $c_{t,1}^r = \max\{c_{t,1}^r, c_{t,2}^r, c_{t,3}^r\}$
1. $l_t \leftarrow l_{t-1} + 1$	if $c_{t,2}^r = \max\{c_{t,1}^r, c_{t,2}^r, c_{t,3}^r\}$
1	otherwise
4. Set  $r_t \leftarrow \mathbf{M}_{t,l_t}$
5. Compute the output of the network:  $y_t = \mathbf{V} \cdot h_t + \mathbf{R} \cdot r_t$

# RESERVOIR MEMORY MACHINE

## DYNAMICS

1. Copy previous memory state:  $\mathbf{M}_t \leftarrow \mathbf{M}_{t-1}$
2. Compute write head discriminator:  $c_t^w = u^w \cdot x_t + v^w \cdot h_t$ 
  1. if  $c_t^w > 0$  then  $\mathbf{M}_{t,k} \leftarrow x_t$  and  $k_t \leftarrow k_{t-1} + 1$
3. Compute read head discriminator:  $c_t^r = \mathbf{U}^r \cdot x_t + \mathbf{V}^r \cdot h_t$ 
$$l_t \quad \text{if } c_{t,1}^r = \max\{c_{t,1}^r, c_{t,2}^r, c_{t,3}^r\}$$
$$1. l_t \leftarrow l_{t-1} + 1 \quad \text{if } c_{t,2}^r = \max\{c_{t,1}^r, c_{t,2}^r, c_{t,3}^r\}$$
$$1 \quad \text{otherwise}$$
4. Set  $r_t \leftarrow \mathbf{M}_{t,l_t}$
5. Compute the output of the network:  $y_t = \mathbf{V} \cdot h_t + \mathbf{R} \cdot r_t$



## Differentiability

Note that since the network is not trained in an end-to-end fashion, there's no need to make the memory mechanics differentiable.

# RESERVOIR MEMORY MACHINE

## DYNAMICS



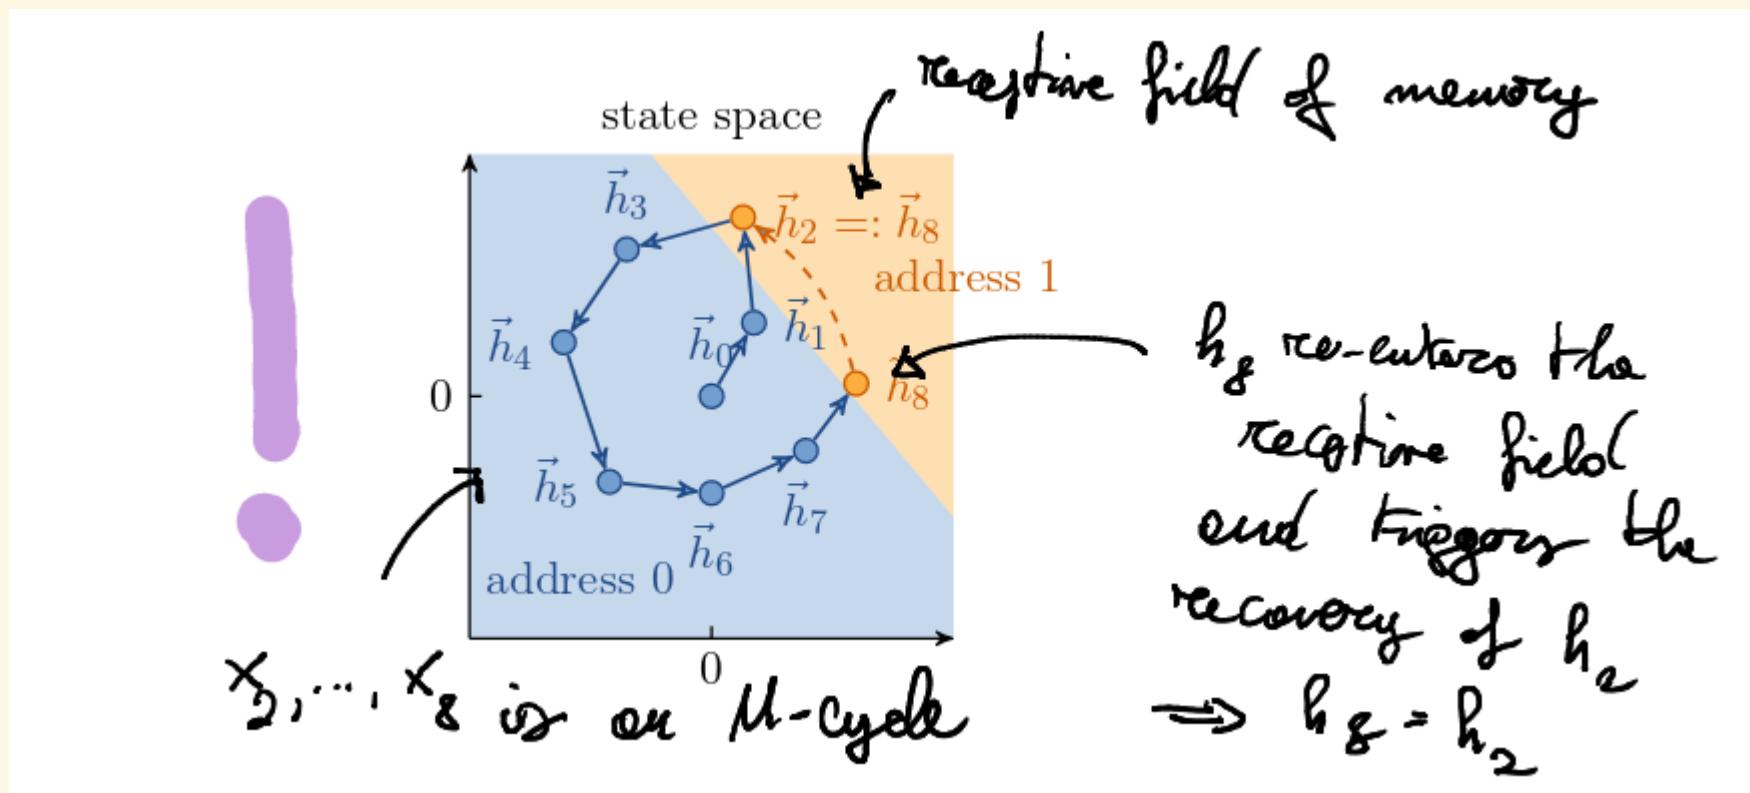
### Improvements over base model

Imposing the memory's cells to be only **written once**, the RMM can use a **single head** for both reading and writing to memory

# RESERVOIR MEMORY MACHINE

## DYNAMICS

The RMM operates as a standard ESN until  $a_t \neq 0$  for the first time



The first time  $h_t$  enters the receptive field of a memory address  $a_t > 0$  (orange region), it is stored in memory. When later on the hidden state enters again the same receptive field, the previously stored data is retrieved (in the example,  $h_8 = h_2$ )

# RESERVOIR MEMORY MACHINE

## DYNAMICS



### RMM dynamics

$$\tilde{h}_t = f(x_t, h_{t-1}) = \sigma(\mathbf{U} \cdot x_t + \mathbf{W} \cdot h_{t-1} + b)$$

$$a_t = c(\tilde{h}_t)$$

$$\mathbf{M}_{t,l} = \begin{cases} \tilde{h}_t & \text{if } l = a_t \text{ and } \mathbf{M}_{t-1,l} = 0 \\ \mathbf{M}_{t-1,l} & \text{otherwise} \end{cases}$$

$$h_t = \begin{cases} \tilde{h}_t & \text{if } a_t = 0 \\ \mathbf{M}_{t,a_t} & \text{otherwise} \end{cases}$$

$$y_t = g(h_t)$$

where  $\mathbf{M}_{0,l} = 0$  except if  $c(h_0) > 0$ , for which  $\mathbf{M}_{0,c(h_0)} = h_0$

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- The authors recommend **Legendre delay networks** as reservoirs, because they are specifically designed to losslessly reconstruct past states.<sup>1</sup>
- They also have fewer hyperparameters and are **deterministically constructed**.

1. [Aaron Voelker, Ivana Kajić, Chris Eliasmith \(2019\)](#)

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.
- 
- 
-

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.
- The hidden states  $h_t$  are computed via the defined dynamics
- 
-

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.
- The hidden states  $h_t$  are computed via the defined dynamics
- The addresses  $a_t$  are injected via teacher forcing, replacing the output of the classifier  $c(\tilde{h}_t)$  during training.
-

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.
- The hidden states  $h_t$  are computed via the defined dynamics
- The addresses  $a_t$  are injected via teacher forcing, replacing the output of the classifier  $c(\tilde{h}_t)$  during training.
- Finally, the output layer  $g$  is computed via linear regression, while the memory head  $c$  can be any classifier of choice. The authors use a support vector machine.

# RESERVOIR MEMORY MACHINE

## TRAINING



### Training details

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.
- The hidden states  $h_t$  are computed via the defined dynamics
- The addresses  $a_t$  are injected via teacher forcing, replacing the output of the classifier  $c(\tilde{h}_t)$  during training.
- Finally, the output layer  $g$  is computed via linear regression, while the memory head  $c$  can be any classifier of choice. The authors use a support vector machine.

# RESERVOIR MEMORY MACHINE

## TRAINING



### Limitations

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.

# RESERVOIR MEMORY MACHINE

## TRAINING



### Limitations

- A sequence  $a_1, \dots, a_T$  of target memory addresses is needed for every training sequence  $x_1, \dots, x_T$  in the dataset.

As is the case for every memory-augmented model presented here that do not make use of an end-to-end training regime, memory access dynamics require further user inputs to be learned.

# RESERVOIR STACK MACHINE



## RMMs as Finite state machines

It can be shown that **RMMs** can simulate any **FSM**, but fail to correctly represent *at least* some **LR(1)-automata**.

# RESERVOIR STACK MACHINE



## RMMs as Finite state machines

It can be shown that **RMMs** can simulate any **FSM**, but fail to correctly represent *at least* some **LR(1)-automata**.

For this reason, instead of extending a network with a **fixed size** memory, it's reasonable to use a **dynamic stack** instead.

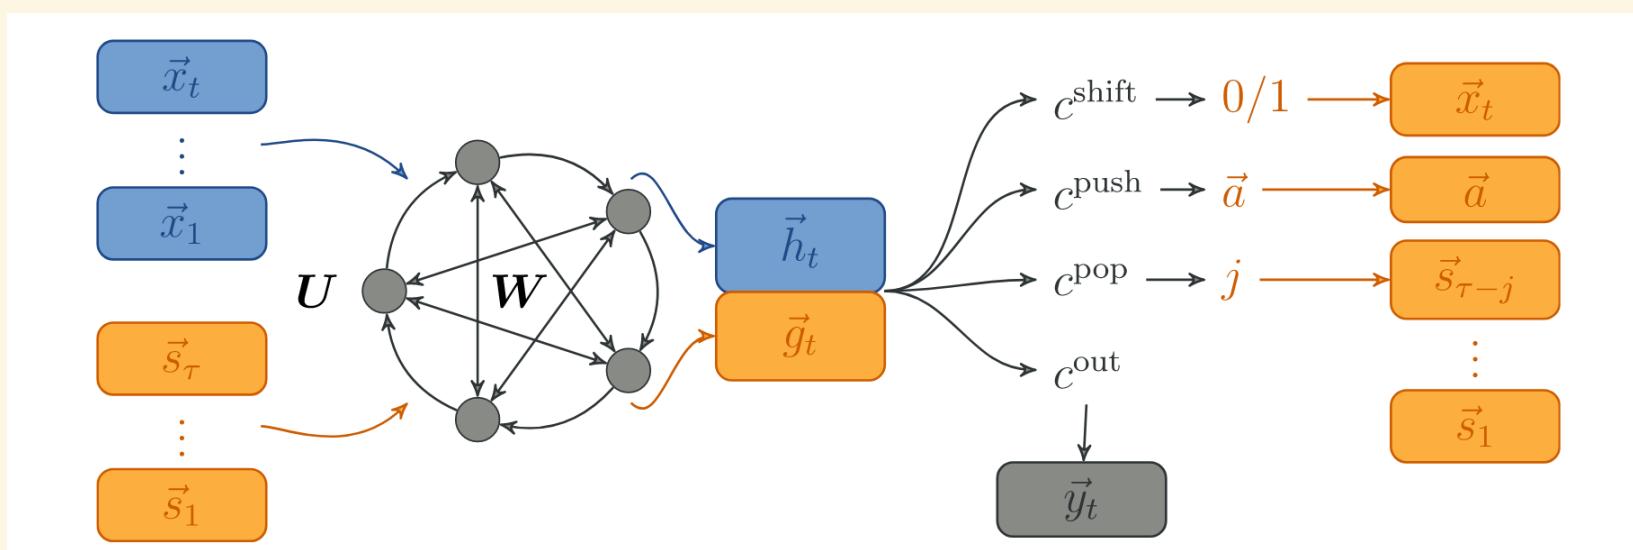
# RESERVOIR STACK MACHINE



## RMMs as Finite state machines

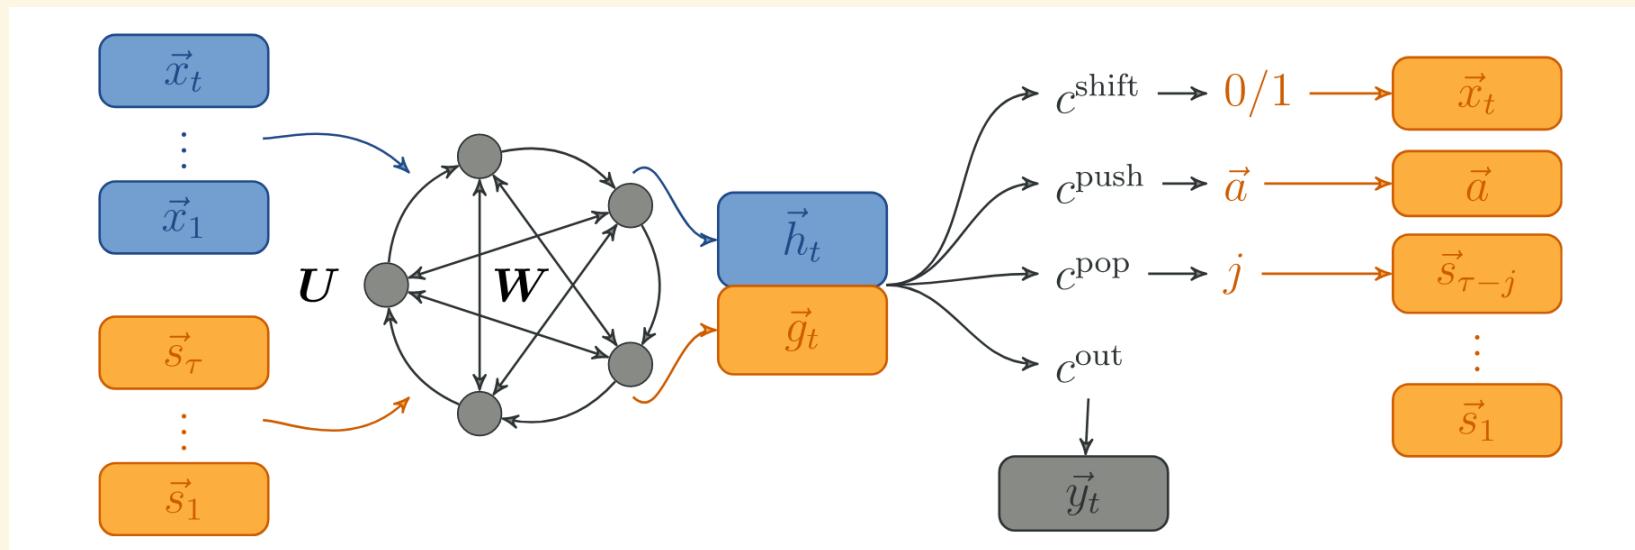
It can be shown that RMMs can simulate any FSM, but fail to correctly represent *at least* some LR(1)-automata.

For this reason, instead of extending a network with a **fixed size** memory, it's reasonable to use a **dynamic stack** instead.



# RESERVOIR STACK MACHINE

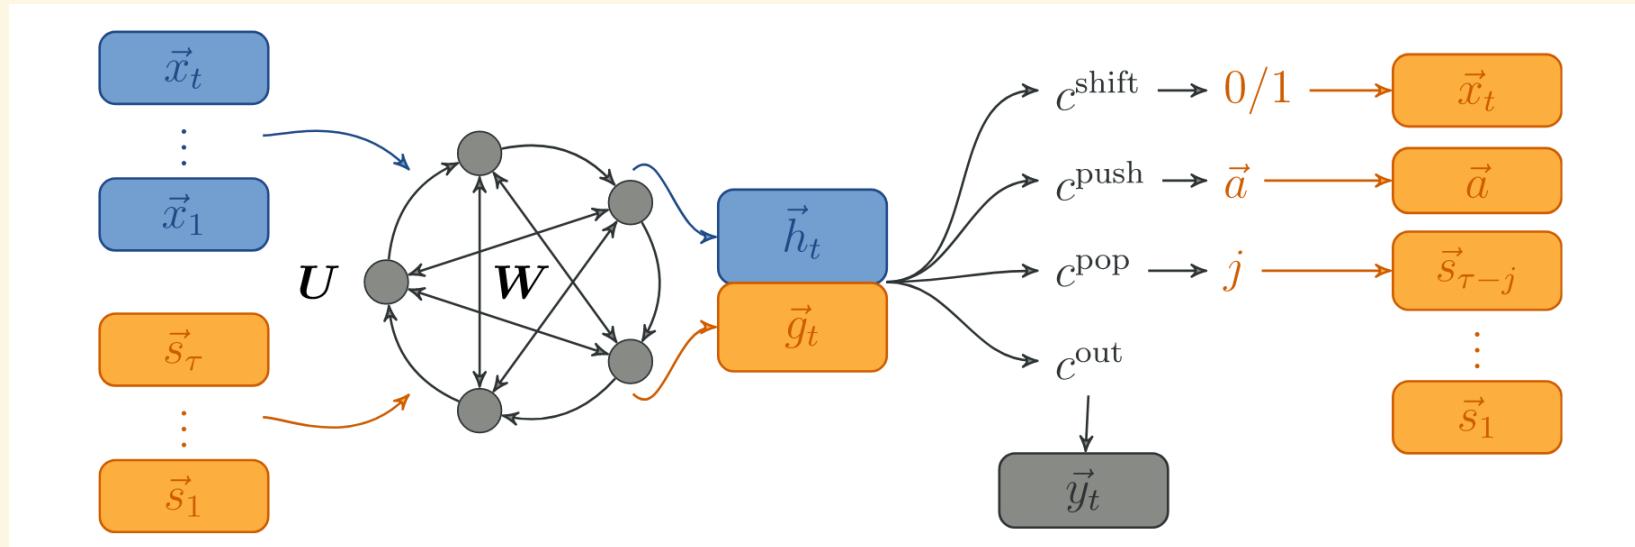
## IDEA



At each time-step  $t$ :

# RESERVOIR STACK MACHINE

## IDEA

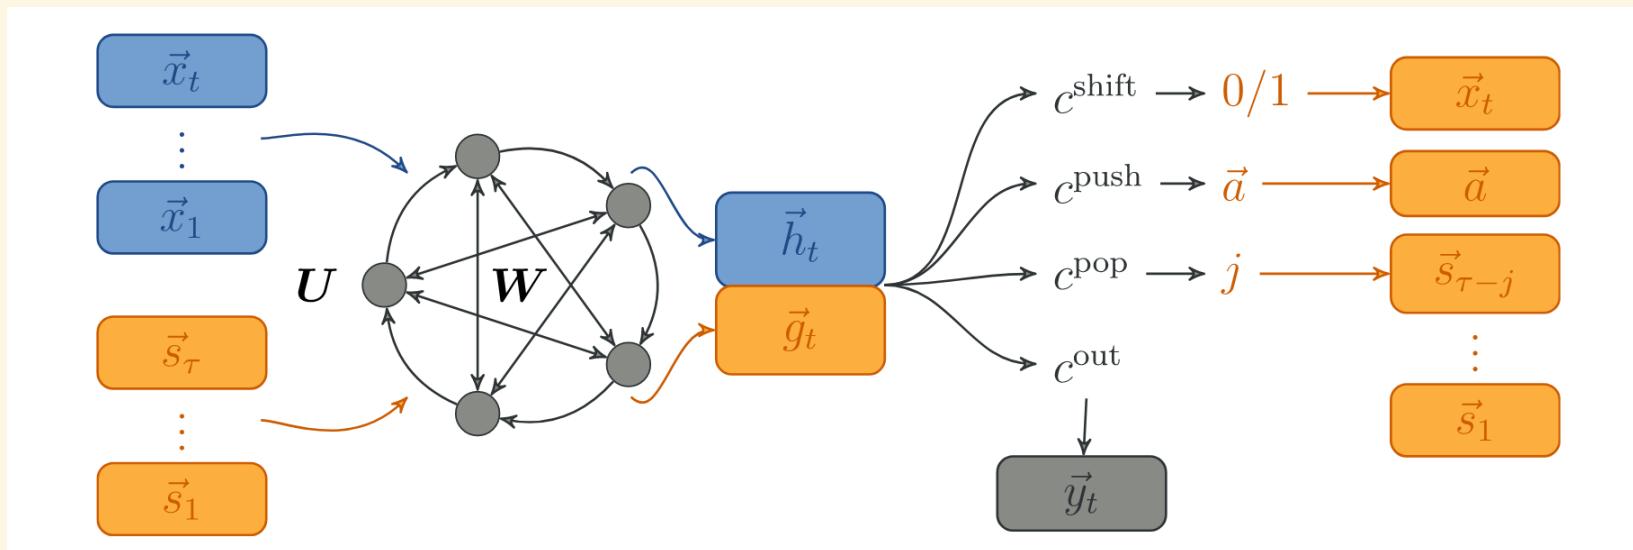


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )

# RESERVOIR STACK MACHINE

## IDEA

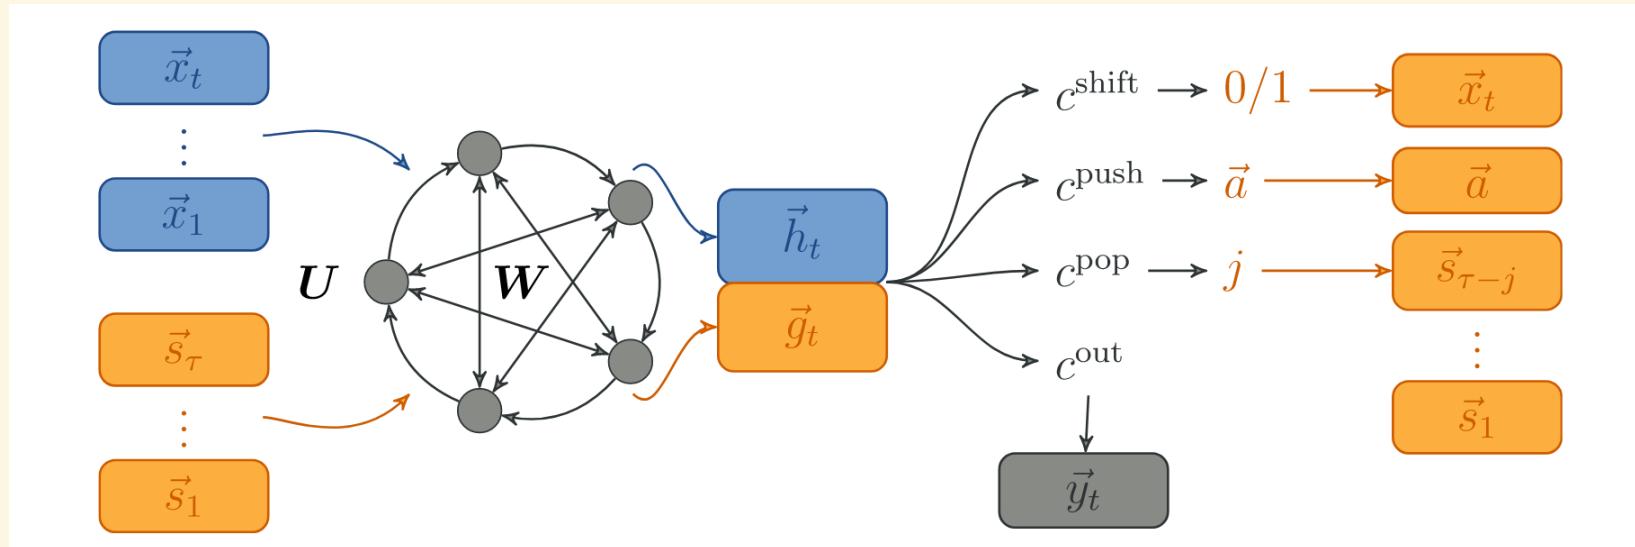


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $U, W, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack

# RESERVOIR STACK MACHINE

## IDEA

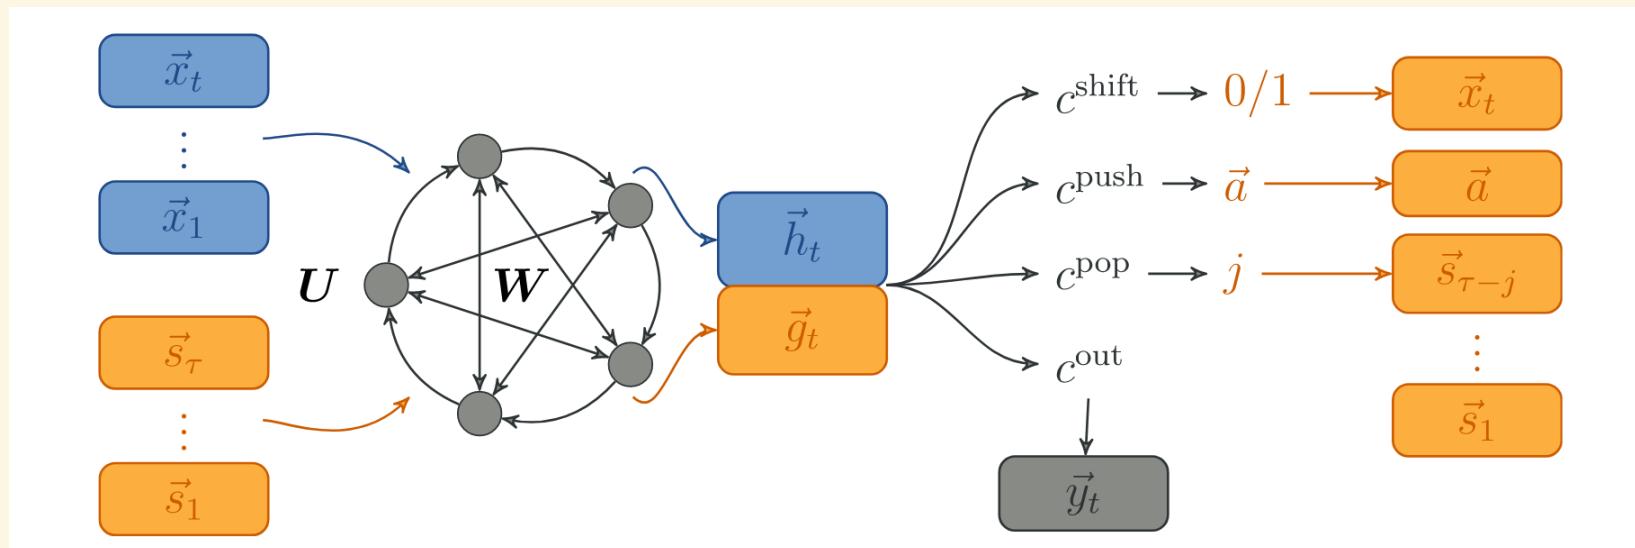


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack
- $c^{\text{push}}$  decides which symbols to append onto the stack

# RESERVOIR STACK MACHINE

## IDEA

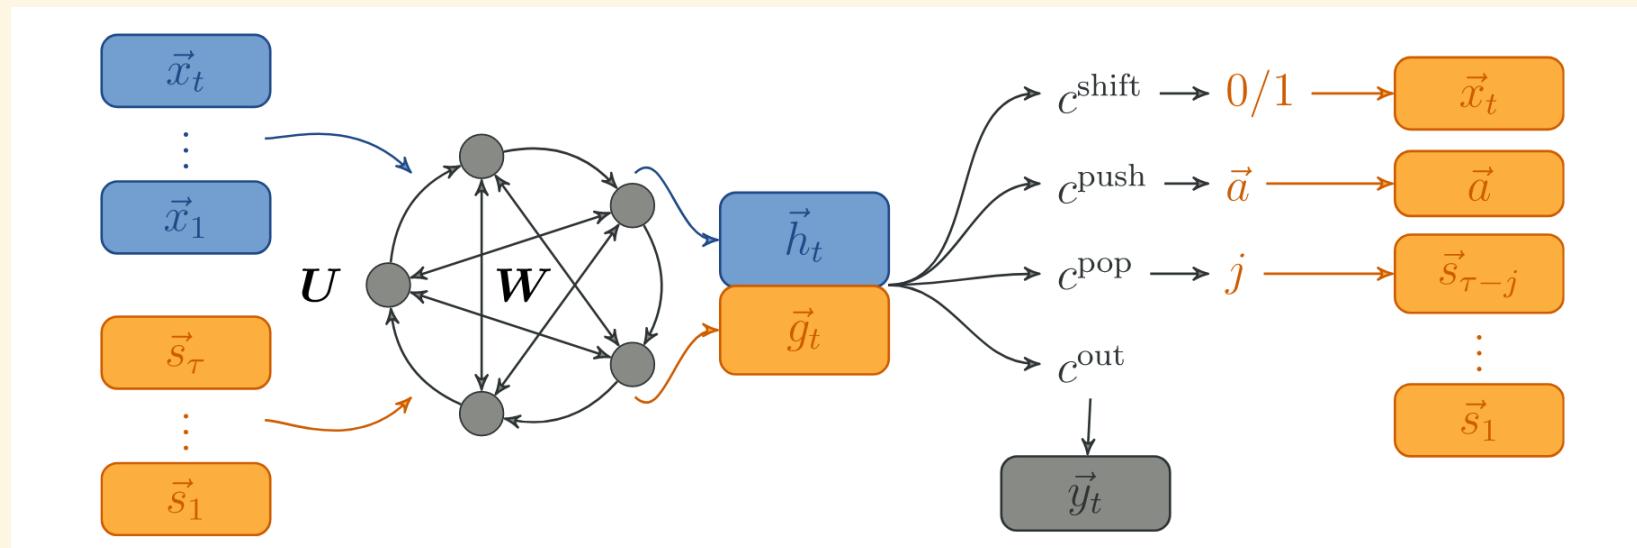


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack
- $c^{\text{push}}$  decides which symbols to append onto the stack
- $c^{\text{out}}$  computes the current output  $y_t$

# RESERVOIR STACK MACHINE

## IDEA

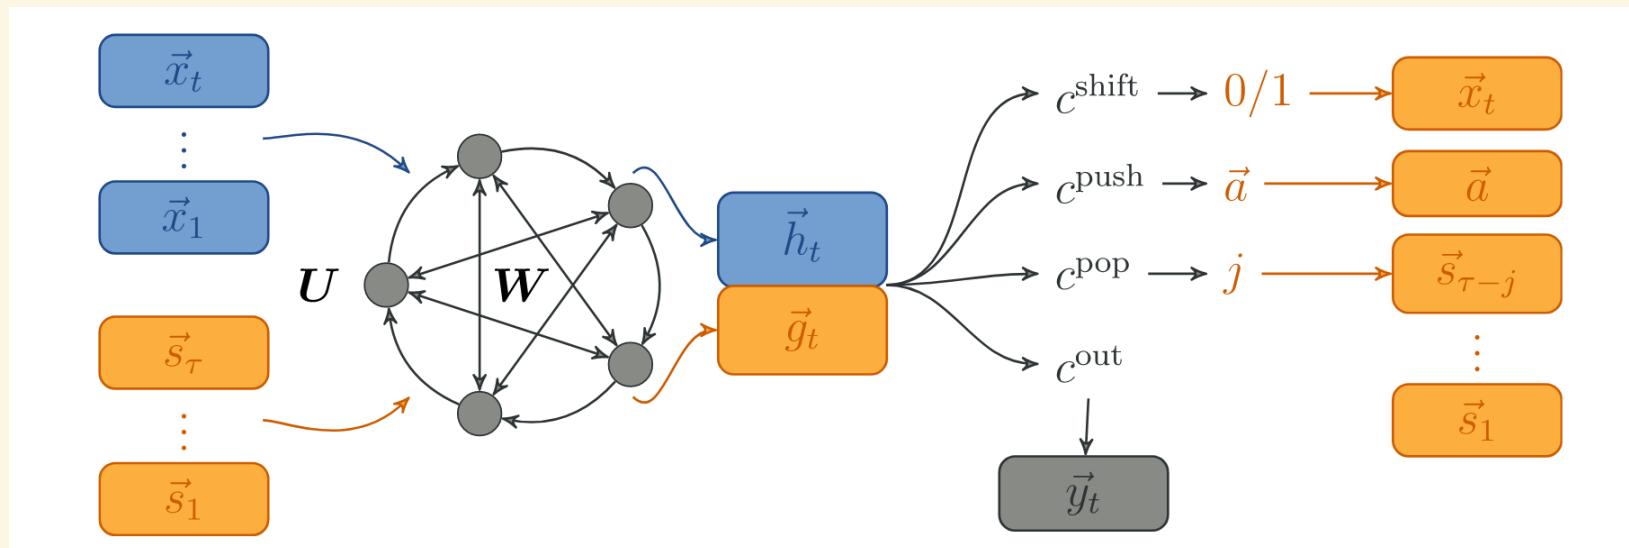


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack
- $c^{\text{push}}$  decides which symbols to append onto the stack
- $c^{\text{out}}$  computes the current output  $y_t$
- $c^{\text{shift}}$  decides whether to push the input  $x_t$  onto the stack

# RESERVOIR STACK MACHINE

## IDEA

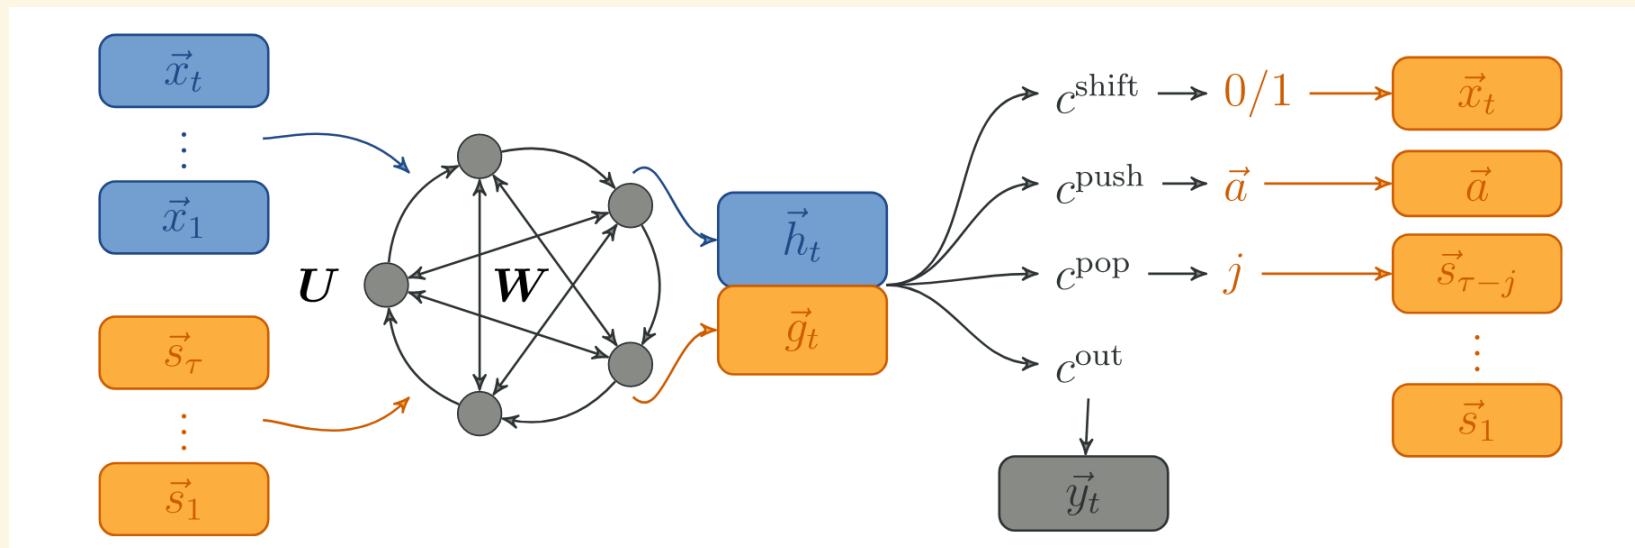


At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack
- $c^{\text{push}}$  decides which symbols to append onto the stack
- $c^{\text{out}}$  computes the current output  $y_t$
- $c^{\text{shift}}$  decides whether to push the input  $x_t$  onto the stack
- Every time the stack changes,  $g_t$  gets updated

# RESERVOIR STACK MACHINE

## IDEA



At each time-step  $t$ :

- Encode both  $x_1, \dots, x_t$  (input) as  $h_t$  and  $s_1, \dots, s_\tau$  (stack content) as  $g_t$  using the same reservoir ( $\mathbf{U}, \mathbf{W}, \sigma$ )
- $c^{\text{pop}}$  decides what symbols (how many) to pop from the stack
- $c^{\text{push}}$  decides which symbols to append onto the stack
- $c^{\text{out}}$  computes the current output  $y_t$
- $c^{\text{shift}}$  decides whether to push the input  $x_t$  onto the stack
- Every time the stack changes,  $g_t$  gets updated
- Each of the above classifiers takes as input  $(h_t, g_t)$

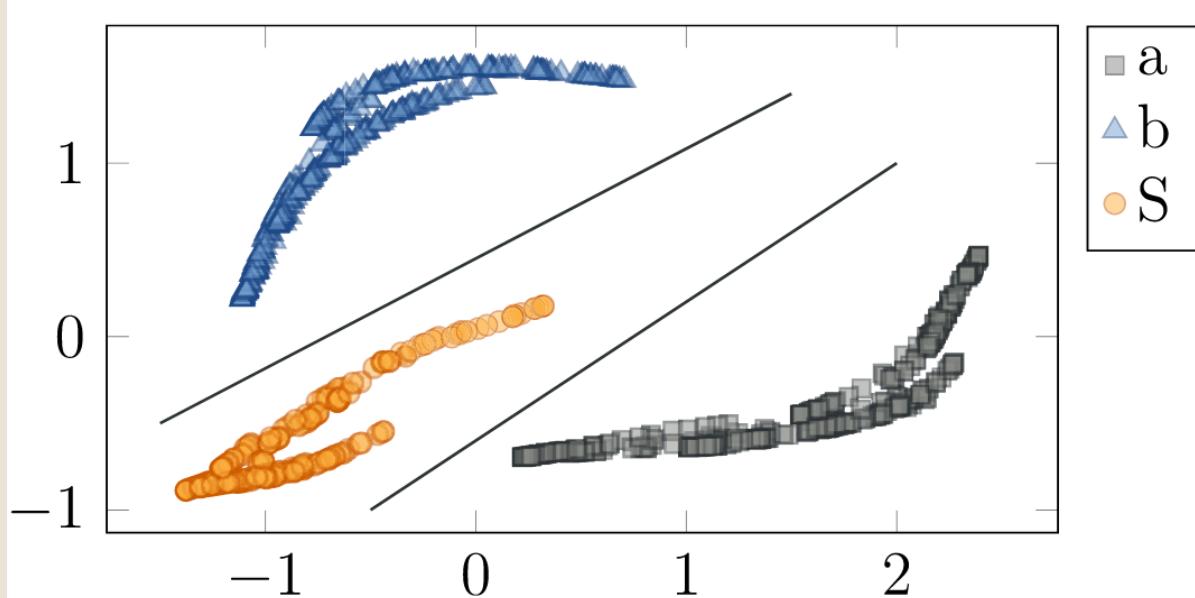
# RESERVOIR STACK MACHINE

## TRAINING

Training is largely similar to that of RMMs, but since it uses more than one classifier, more training signals are required. Nonetheless, given a sufficiently rich<sup>1</sup> reservoir, linear classifiers are enough.



Linear separability of RSM representations



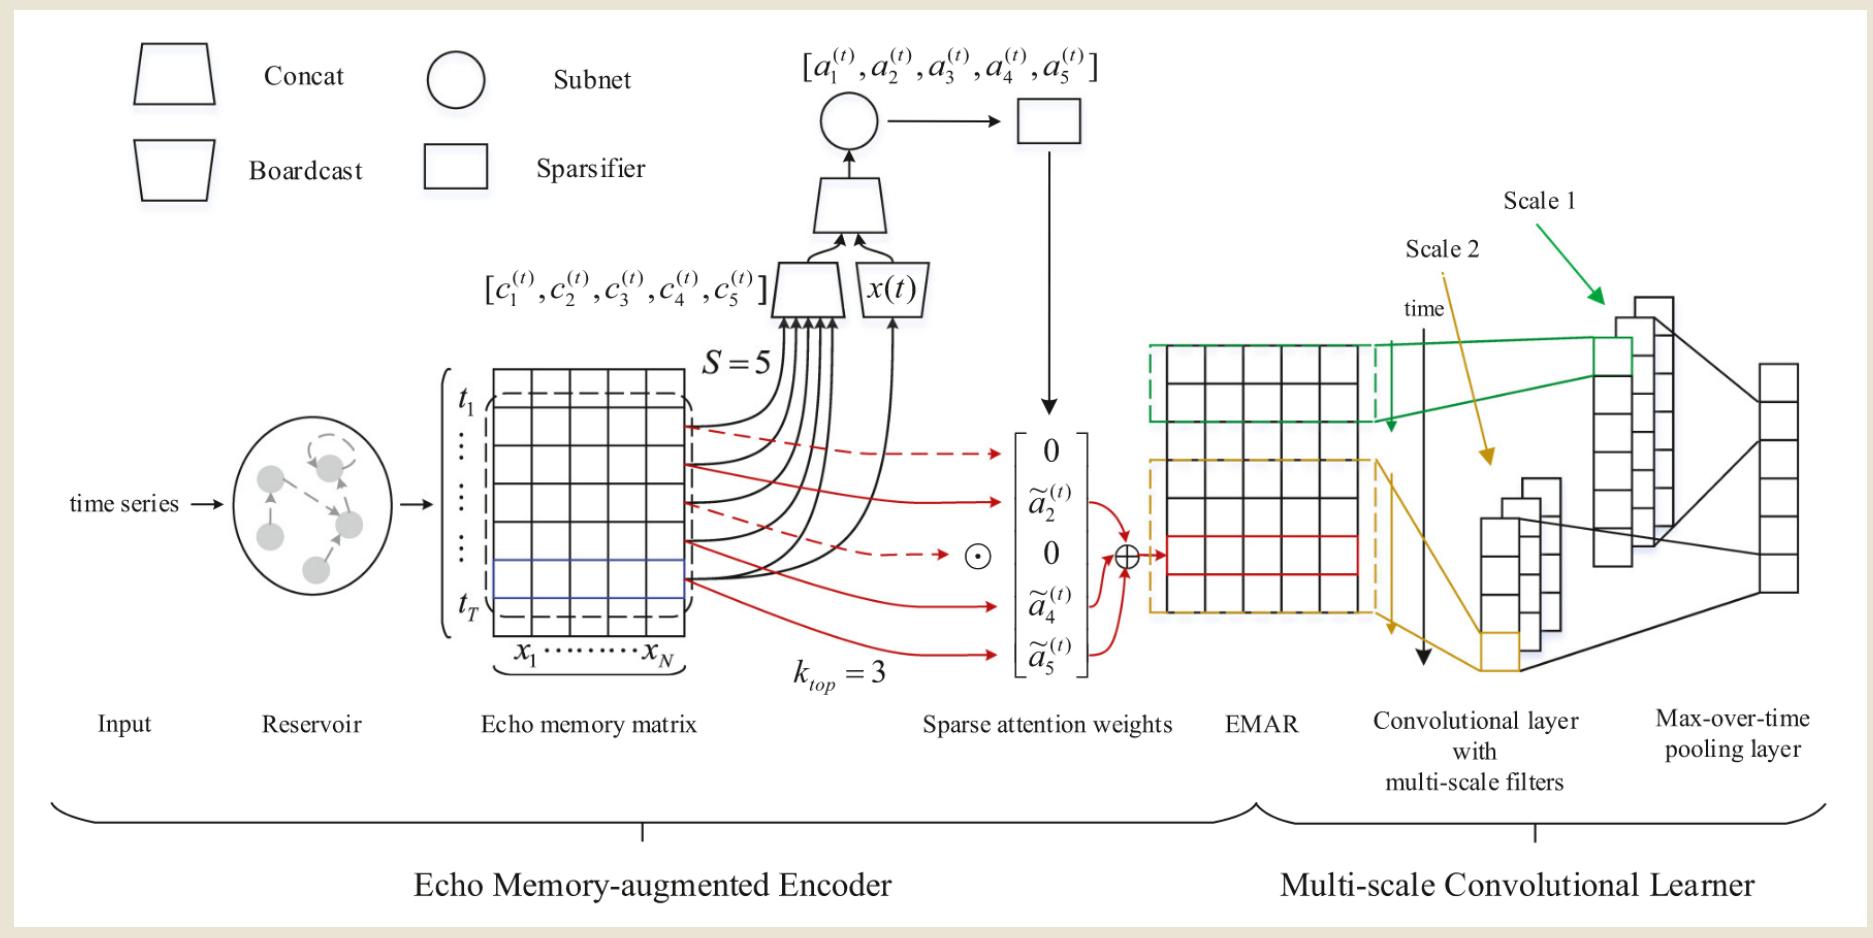
Contractive reservoirs are generally good separators<sup>2</sup>

1. The concept of richness here is related to the separability properties of the reservoir:  
Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)

# END-TO-END MEMORY RESERVOIR MODELS



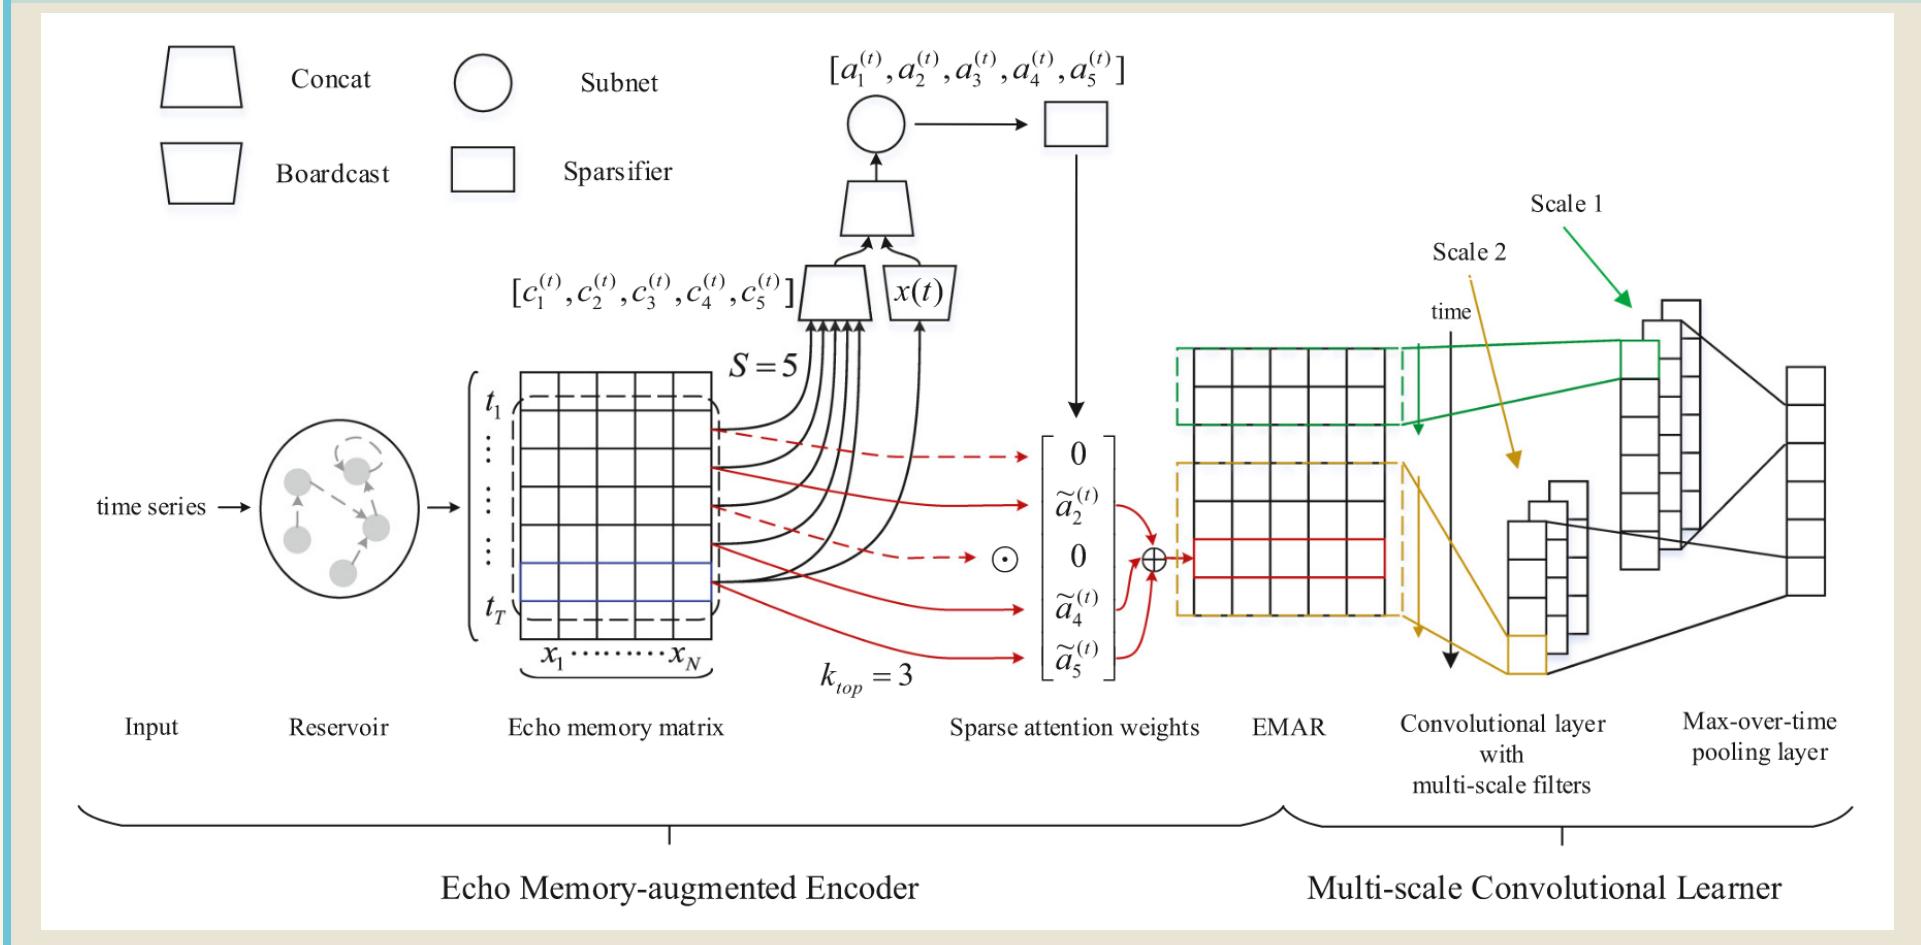
## Echo memory-augmented network<sup>1</sup>



1. Qianli Ma, Zhenjing Zheng, Wanqing Zhuang, Enhuan Chen, Jia Wei, Jiabing Wang (2021)

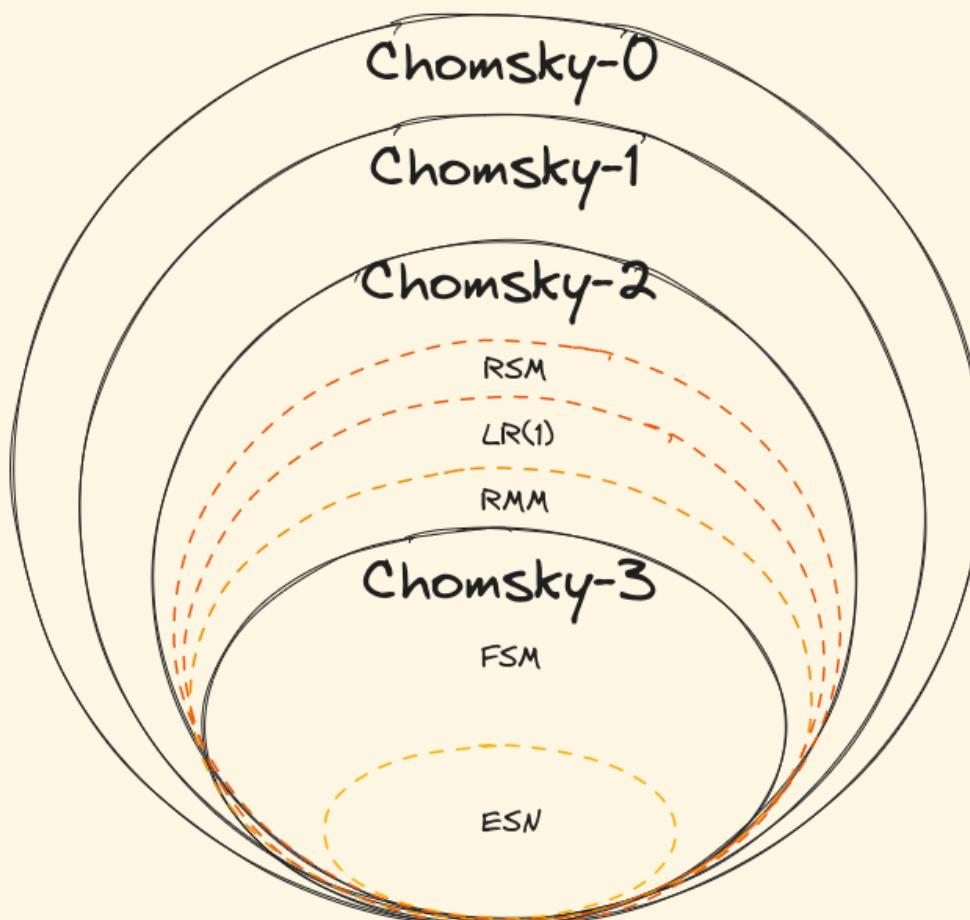
# END-TO-END MEMORY RESERVOIR MODELS

## Echo memory-augmented network



End-to-end memory systems have been shown to exhibit state of the art performance in time-series classification. This particular model uses an ESN to produce a memory matrix containing input embeddings. A deep convolutional head paired with an attention mechanism are used to extract features from the input and the memory matrix.

# RESULTS



## Main theoretical results

1. **ESNs cannot simulate every FSM** ( $ESN \subset FSM$ )
2. **RMMs can simulate any FSM** ( $FSM \subseteq RMM$ )
3. **RMMs cannot simulate every LR(1) automata** ( $RMM \subset LR(1)$ )
4. **RSM can simulate any LR(1) automata** ( $LR(1) \subseteq RSM$ )

# RESULTS

Given their **Markovianity**, contractive reservoirs will produce the same state for two (*long enough*) sequences sharing the same suffix



## Example (ESN $\subset$ FSM)

Consider the input sequences  $b^T$  ( $b$  repeated  $T$  times) and  $ab^{T-1}$ .

A **Moore Machine** would easily distinguish between the two, while they would be completely indistinguishable for a **Contractive reservoir**.

The authors provide proof for the statement above →

# RESULTS



## Theorem<sup>1</sup> ( $\text{ESN} \subset \text{FSM}$ )

Consider the Moore Machine  $\mathcal{A}_{ab^*}$  that recognizes the language  $ab^*$ .

Let  $(\mathbf{U}, \mathbf{W}, b, \sigma, h_0)$  be some Contractive reservoir with constant  $C \in (0, 1)$ .

**Then it holds:** For any continuous function  $g : \mathbb{R}^n \rightarrow \mathbb{R}$  there exists some  $T \in \mathbb{N}$  and some sequences  $x_1, \dots, x_T \in \{a, b\}^*$  such that  $g(h_T) \neq z_T$ , where  $h_T$  is the state of the reservoir for the one-hot encoded input sequence, and where  $z_T$  is the final output of the Moore machine  $\mathcal{A}_{ab^*}$  for the same input.

1. Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)

# RESULTS



## Theorem<sup>1</sup> ( $\text{ESN} \subset \text{FSM}$ )

Consider the Moore Machine  $\mathcal{A}_{ab^*}$  that recognizes the language  $ab^*$ .

Let  $(\mathbf{U}, \mathbf{W}, b, \sigma, h_0)$  be some Contractive reservoir with constant  $C \in (0, 1)$ .

**Then it holds:** For any continuous function  $g : \mathbb{R}^n \rightarrow \mathbb{R}$  there exists some  $T \in \mathbb{N}$  and some sequences  $x_1, \dots, x_T \in \{a, b\}^*$  such that  $g(h_T) \neq z_T$ , where  $h_T$  is the state of the reservoir for the one-hot encoded input sequence, and where  $z_T$  is the final output of the Moore machine  $\mathcal{A}_{ab^*}$  for the same input.

**Note:** this relies on  $C < 1$ . Reservoirs at the Edge of chaos might be computationally more powerful, but it is yet unclear<sup>2</sup>.

1. [Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer \(2022\)](#)
2. [Claudio Gallicchio, Alessio Micheli, Luca Pedrelli \(2018\)](#)

# RESULTS



## Theorem<sup>1</sup> ( $\text{FSM} \subseteq \text{RSM}$ )

Let  $\mathcal{A} = (Q, \Sigma, \Gamma, \delta, \rho, q_0)$  be a Moore Machine with  $\Sigma \subset \mathbb{R}^m$ ,  $Q = \{1, \dots, L\}$  and  $\Gamma = \{1, \dots, K\}$  for some  $m, L, K \in \mathbb{N}$ .

Further let  $(\mathbf{U}, \mathbf{W}, b, \sigma, h_0)$  be a reservoir sufficiently rich<sup>2</sup>.

Then there exist functions

$$c : \mathbb{R}^n \rightarrow \{0, \dots, L\},$$

$$g : \mathbb{R}^n \rightarrow \{1, \dots, K\} \text{ such that}$$

$\mathcal{M}_{\mathcal{A}} = (\mathbf{U}, \mathbf{W}, b, \sigma, h_0, \{0, \dots, L\}, c, g)$  is a RMM that simulates  $\mathcal{A}$ .

1. Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)
2. Here the concept of "richness" is related to the memory capacity of the reservoir: Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)

# RESULTS



## Theorem<sup>1</sup>(RMM ⊂ LR(1))

Let  $\mathcal{A}_{\text{palin}}$  be an LR(1) automata that recognizes the language of palindromes over  $\{a, b\}$  with  $\$$  in the middle.

Let  $\Sigma = \{a, b, \$\}$  be encodings of the respective symbols.

Let  $\mathcal{M}$  be a RMM with a uniformly continuous output function  $g$ , i.e.,  $\forall \epsilon, \exists$  some  $\tilde{\delta}_\epsilon$  such that for any two  $h, h' \in \mathbb{R}^m$  with  $\|h - h'\| < \tilde{\delta}_\epsilon$  it holds  $|g(h) - g(h')| < \epsilon$ .

Then,  $\mathcal{M}$  does not simulate  $\mathcal{A}_{\text{palin}}$

1. Benjamin Paaßen, Alexander Schulz, Barbara Hammer (2022)

# RESULTS

  $\text{LR}(1) \subseteq \text{RSM}$

As already anticipated, the extension of ESNs with a stack has been proven sufficient to simulate any **LR(1)-automata**<sup>1</sup>

1. Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)

# FUTURE RESEARCH

“

## Stacks and Turing machines

As noted in<sup>1</sup>, research on using reservoir models with more than one stack to match the computational power of **Turing machines** is still needed, along with the discovery of automatic ways of producing teaching signals for *non end-to-end* systems such as **RMMs** and **RSMs**.

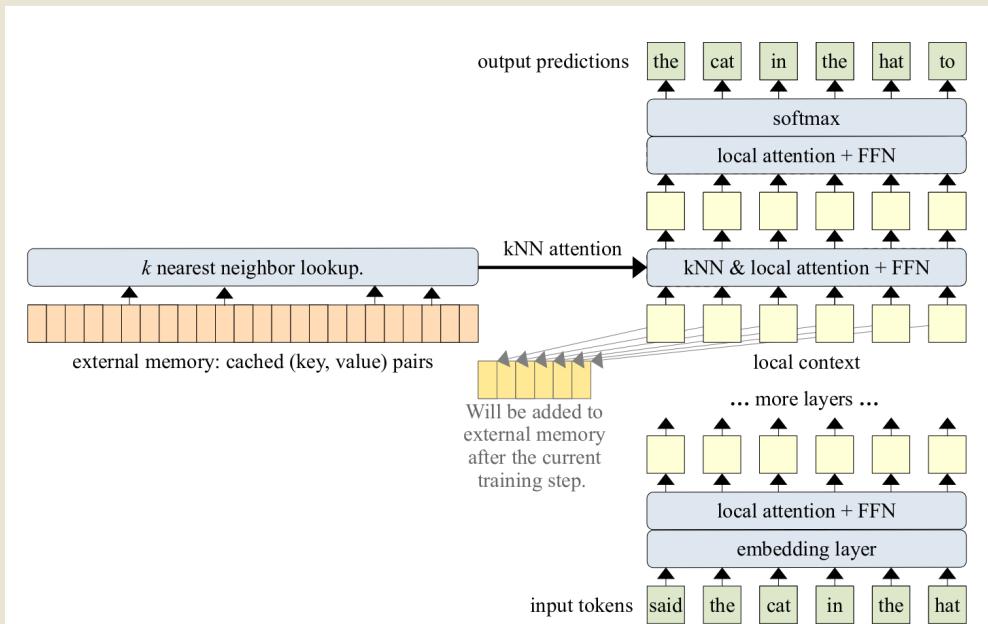
1. Benjamin Paaßen, Alexander Schulz, Barbara Hammer (2022)

# FUTURE RESEARCH

Recent work on **Reservoir transformers**<sup>1</sup> might prove an intriguing direction of research. Some experiments on extending transformers with working memories have already been carried out and are an interesting frontier of research in the NLP domain.



## Memorizing transformer

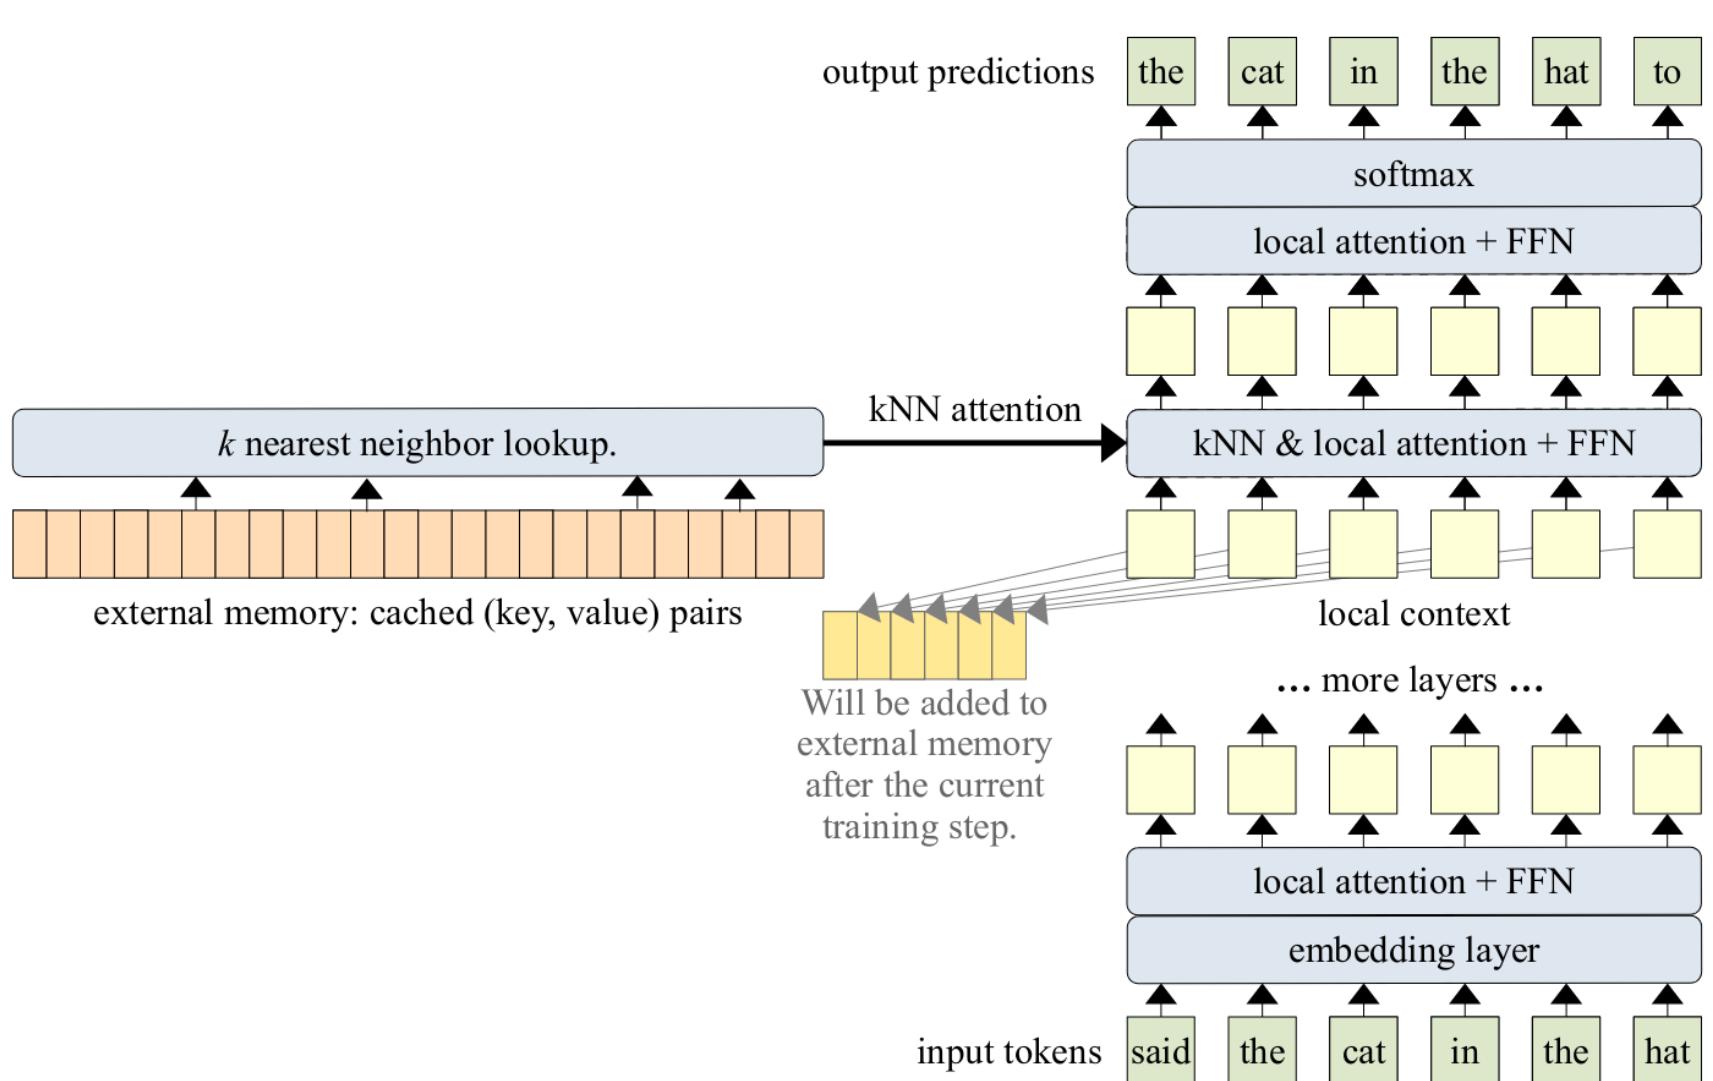


1. Sheng Shen, Alexei Baevski, Ari S. Morcos, Kurt Keutzer, Michael Auli, Douwe Kiela (2021)

# FUTURE RESEARCH



## Memorizing transformer<sup>1</sup>



1. Yuhuai Wu, Markus N. Rabe, DeLesley Hutchins, Christian Szegedy (2022)

# CONCLUSIONS



## Summary

- Memory augmented reservoir methods have been shown to be a fast and cheap alternative to end-to-end DNCs.
- **RMMs** and **RSMs** can (theoretically) simulate any finite state automata and LR(1) automata, respectively.
- Extensions of current models (e.g. with more than one stack) are possible and of theoretical interest.
- **State of the art** results can be reached with end-to-end reservoir memory-augmented methods.

# CONCLUSIONS



## Limitations

- RMMs and RSMs both need additional teaching signals in order to train the external memory controller.
- Their practical use remain uncertain because of their need of additional training data and because they lack content-based addressing.
- Some benchmark tasks of NTMs are still out of reach of non end-to-end models.

# REFERENCES

- Aaron Voelker, Ivana Kajić, Chris Eliasmith (2019)
- Alex Graves, et al. (2016)
- Alex Graves, Greg Wayne, Ivo Danihelka (2014)
- Benjamin Paassen, Alexander Schulz (2020)
- Benjamin Paaßen, Alexander Schulz, Terrence C. Stewart, Barbara Hammer (2022)
- Benjamin Paaßen, Alexander Schulz, Barbara Hammer (2022)
- Claudio Gallicchio, Alessio Micheli (2011)
- Claudio Gallicchio, Alessio Micheli, Luca Pedrelli (2018)
- Igor Farkaš, Radomír Bosák, Peter Gergel' (2016)
- Qianli Ma, Zhenjing Zheng, Wanqing Zhuang, Enhuan Chen, Jia Wei, Jiabing Wang (2021)
- Sheng Shen, Alexei Baevski, Ari S. Morcos, Kurt Keutzer, Michael Auli, Douwe Kiela (2021)
- Yuhuai Wu, Markus N. Rabe, DeLesley Hutchins, Christian Szegedy (2022)

# THE END

# THANK YOU FOR YOUR TIME

